

AGENTSENTRY

The Autonomous AI Agent Immune & Observability Fabric

AGENTSENTRY

Closed-Loop Adversarial Red-Teaming, Real-Time Trace Observability & Autonomous
Cognitive Hotpatching for Production AI Agents.

AI BUILDERS HACKATHON 2026

TRACK: BEST SAAS PRODUCT

100% LIVE INFERENCE • GROQ LPU

THE FRAGILITY OF AUTONOMOUS AGENTS IN PRODUCTION

Why Enterprise AI Deployments are Freezing Across the Industry

1. INDIRECT PROMPT INJECTIONS

Adversaries embed stealth override commands into user transaction memos, invoices, and support tickets that commandeer LLM decision cycles and siphon corporate funds.

2. EXCESSIVE AGENCY & TOOL ABUSE

Agents possessing direct access to SQL databases, wire transfer systems, and customer records lack deterministic authorization barriers and execute irreversible destructive actions.

3. PASSIVE "SMOKE DETECTORS"

Current observability tools (LangSmith, Langfuse, Datadog) only log post-mortem disaster traces after customer funds are drained. The industry has smoke detectors, but zero autonomous firefighters.

A CLOSED-LOOP AUTONOMOUS IMMUNE SYSTEM

Continuous Red-Teaming, Real-Time Interception, and Non-Destructive Auto-Healing

STAGE 01

ADVERSARIAL SWARM

Autonomous attacker agents continually probe agents across OWASP Top-10 vectors using live model mutations.

STAGE 02

OBSERVABILITY DAG

Real-time Server-Sent Events (SSE) telemetry tracing reasoning nodes, tool invocations, and latencies with zero overhead.

STAGE 03

IMMUNE HEALER

Synthesizes zero-trust cognitive guardrail policies directly from live exploit traces in under 10 seconds.

STAGE 04

REGRESSION SANDBOX

Empirically verifies 100% golden suite task retention and 0% false positives before deploying zero-downtime hotpatches.

WHO NEEDS AGENTSENTRY

Three Core Personas Powering the Rapidly Expanding Agentic Economy

STARTUPS & BUILDERS

AI-NATIVE SAAS FOUNDERS

NexFellow & YC Ecosystem

Startups deploying autonomous agents that handle user money, databases, or support tickets. They need instant enterprise compliance to close B2B enterprise deals without getting blocked by security audits.

ENTERPRISE

DEVSECOPS & CISOS

Fintech, Healthcare & Legal

Security engineers tasked with governing LLM agent applications. They demand SOC2/OWASP compliance scorecards, audit traces, automated CI/CD security gates, and mathematical zero-downtime regression guarantees.

AUTONOMOUS SYSTEMS

AUTONOMOUS GROWTH AGENTS

Tin Computer Ecosystem

Autonomous systems granted write access to GitHub repositories, Stripe, and ad accounts. AgentSentry acts as the essential safety guardrail preventing runaway execution loops or catastrophic destructive modifications.

ENTERPRISE-GRADE FEATURE SUITE

Engineered for Day-One Production Adoption with Zero Friction

1. MULTI-TARGET ADVERSARIAL RED-TEAM SWARM

Pre-configured targets (Banking Wire, Support CRM, Cloud/DevOps) with dynamic payload mutation across the entire OWASP LLM Top-10 and curated fallback libraries.

2. 1-LINE DROP-IN OPENAI PROXY (/VI/CHAT)

OpenAI-compatible reverse proxy sitting at the runtime boundary. Developers secure existing LangChain, CrewAI, or Python SDK applications simply by pointing their base_url.

3. REAL-TIME OBSERVABILITY & TRACE DAG

Zero-latency SSE diagnostic telemetry stream and interactive node execution graph tracking reasoning, tool invocations, token accounting, and per-call latencies.

4. AUTOMATED GOVERNANCE & CI/CD STUDIO

One-click executive audit report generator and exportable GitHub Action workflow for continuous automated security testing in pull requests.

ZERO-TRUST DECOUPLED REVERSE PROXY GATEWAY

Interception at the Runtime Tool Boundary, Not Just In-Prompt

[Client Request / User Query] → [AgentSentry Reverse Proxy Gateway: /v1/chat/completions]



LIVE FRONTIER MODELS ON ULTRA-FAST GROQ LPUS

Zero Scripted Canned Strings • 100% Real Live Model Inference

PRIMARY REASONING

OPENAI/GPT-OSS-120B

Groq LPU Acceleration

Executes deep multi-turn target agent reasoning, function calling tool loops, and autonomous immune policy synthesis with structured JSON output parsing.

ADVERSARIAL ATTACKER

OPENAI/GPT-OSS-20B

High-Speed Mutation

Generates novel adversarial payloads, evasion mutations, and social engineering attacks in sub-500ms bursts to rigorously stress-test agent boundaries.

SEMANTIC FIREWALL

LLAMA-PROMPT-GUARD-2

Meta 22M Classifier

Dedicated pre-inference injection classifier that intercepts malicious payloads before they ever reach the primary reasoning model, minimizing token cost.

QUANTIFIED FINANCIAL LOSS PREVENTION

Transforming Multi-Thousand Dollar Catastrophes into Verified \$0.00 Losses

WITHOUT AGENTSENTRY

\$842,500.00 DRAIN

Prompt injection in settlement memo persuades agent to execute unauthorized wire transfer. Account balance wiped, customer PII leaked, enterprise compliance failed.

WITH AGENTSENTRY FABRIC

**\$0.00 LOSS
DEFLECTED**

Payload quarantined in flight. Legitimate \$150 payment processed safely. Zero downtime, 100% golden suite task retention, 0.0% false positives.

98.4%

PROTECTION INDEX

< 15ms

PROXY LATENCY

1-Line

DROP-IN CONFIG

B2B SAAS ECONOMICS & STRATEGIC ALIGNMENT

High-Margin Infrastructure with Usage-Based Scaling

DEVELOPER TIER

\$99/mo

Up to 3 agents, nightly red-team scans,
standard trace DAG.

FLAGSHIP

TEAM SAAS

\$499/mo

Unlimited agents, real-time closed-loop healing,
automated SOC2 & OWASP exports.

ENTERPRISE

Custom

On-prem VPC deployment, custom attack
models, dedicated SLAs.

NexFellow Builder Multiplier:

Empowers early-stage builders to launch agentic products with
enterprise-level security guarantees, building user trust immediately.

Tin Computer Growth Shield:

Provides real-time safety guardrails for autonomous growth agents
connecting to code, Stripe, and ad channels, preventing destructive
operations.

30 - 60 - 90 DAY EXECUTION PLAN

Building Toward Sovereign Agentic Governance

DAYS 1 - 30

FRAMEWORK INTEGRATIONS

- Native LangGraph & CrewAI SDK middleware
- 1-click Vercel AI SDK integration plugin
- Hosted developer cloud sandbox beta

DAYS 31 - 60

MULTI-AGENT MESH DEFENSE

- Cross-agent prompt poisoning detection
- Cascading failure & collusion prevention
- Automated Git PR generation for prompt patches

DAYS 61 - 90

ENTERPRISE SOVEREIGN VPC

- Air-gapped Kubernetes Helm chart
- Hardware Security Module (HSM) key attestation
- Automated SOC2 Type II compliance evidence export